

1 Convolutional Networks

By: Jussi Martin

1.1 Preface

This post is written for the AI Helsinki study group *Image and Video Statistics*. It is based on the Chapter 9 of the book *Deep Learning* by Ian Goodfellow, Yoshua Bengio and Aaron Courville, which is under preparation and available online at <http://www.deeplearningbook.org>.

We will assume the reader to be familiar with basic components of artificial neural networks, such as weights and biases. Some mathematical preliminaries are introduced in the text.

Our goal is to introduce basic concepts of the convolutional neural networks, explain where they are used and what kind of benefits the convolutional structure offers.

1.2 Definition

We begin by introducing the mathematical notion of convolution. Let $x(t)$ and $w(t)$ two functions then we define their convolution $x * w$ as

$$(x * w)(t) = \int x(a)w(t - a)da.$$

In practical applications x could be some signal depending on the time t and w a filter applied to it.

Obviously the data for which the convolution is applied usually is obtained only as a discrete input, in these cases we define the convolution by using summation:

$$(x * w)(t) = \sum_{a=-\infty}^{\infty} x(a)w(t - a).$$

Moreover, the data might be two dimensional, which is the case with images, then we use the following definition:

$$S(i, j) = (I * K)(i, j) = \sum_m \sum_n I(m, n)K(i - m, j - n)$$

where I and K represent two images and i and j are the pixel coordinates.

In practice the input grid is finite, the filter is defined on a smaller grid and usually only applied when its coordinates are fully contained in the input grid. These technical details will be covered soon.

We say that a neural network is convolutional if it uses convolutional filtering in at least one of its layers.

1.3 Use of Convolutional Networks

Used to process data that has a grid-like structure (e.g. images, 1-D: sound?)

1.4 Benefits (submerge with previous section?)

*Modelling of large inputs while being computationally efficient.

*Invariants to small translations of inputs.

1.5 Tensors

We adopt the Deep Learning book's convention of defining tensors as multidimensional arrays of real numbers.

Namely, 0-D tensors are just real numbers, 1-D tensors are arrays

$$T = (T_1, \dots, T_n)$$

of real numbers, 2-D tensors are arrays

$$T = ((T_{1,1}, \dots, T_{1,n}), \dots, (T_{m,1}, \dots, T_{m,n}))$$

of arrays of real numbers (with mutually equal length), and so on.

Basically our tensors can be viewed as D-dimensional grids of real numbers.

Example in 2-D:

$$((1, 2, 3), (4, 5, 6), (7, 8, 9)) \simeq \begin{array}{c|c|c} 1 & 2 & 3 \\ \hline 4 & 5 & 6 \\ \hline 7 & 8 & 9 \end{array}$$

In general case the sizes of the axes do not need to match.

1.6 Convolution via Tensors

*Stride

*Channels

*Formulas, etc.

(convolution with 1×1)... traditional fully connected layer can be seen as other extreme case, with zero stride and kernel grid size been equal to the input grid size.

1.7 Examples (submerge with previous section?)

(Pic of 2-D convolution, should be earlier?)

1.8 Sparse Connectivity

One benefit of using convolutional layers is that there are fewer connections between the nodes. Only as many input nodes as there are weights in the convolution kernel are connected to every single node in the output nodes and vice versa.

Pic: 1-D, connections from inputs to single node in the outputs

Pic: 1-D, connections from single input to nodes in the outputs

This is in drastic contrast to fully connected layers where there are connections from every input node to a given output node and vice versa.

For example, if we have an image with 28×28 pixels grid, like the images in the MNIST data set have, and we convolve it with 3×3 filter. Then there are only $3 \cdot 3 = 9$ connections from input nodes to single output node, where as in fully connected layer there would be $28 \cdot 28 = 784$ connections to single output node. And similarly with roles of the input and output nodes reversed.

1.9 Growing Receptive Fields (optional ?)

1.10 Parameter Sharing

Another benefit of using convolution is that there are fewer parameters needed. Namely the weights of the convolution kernel are shared between similar type of connection.

Pic: illustration of the parameter sharing

Sometimes even the biases are shared in similar manner, but this type of parameter sharing is not often used.

1.11 Zero Padding

One drawback of using convolutional layers is that the layers shrink, which opposes limit to the depth of the network.

Pic: 1-D network with shrinking layers

This however can be compensated by adding zeros to the boundary of each neuron layer. That way we can prevent the sizes of the layers shrinking and obtain deeper networks.

Pic: 1-D network with zero padding and same size layers

The method described above is known as zero padding. One way to do it is to add just enough zeros that the layer sizes stays the same. Other more extreme option is to add so many zeros that the convolution can visit every neuron k times, where k is the size of the convolution kernel.

Pic: 2-D, original layer and layer with minimal padding

Pic: 2-D, original layer and layer with maximal padding

One may notice that zero padding makes the outputs near the boundary to have less information from the inputs, especially if there is maximal amount of padding used. This may be compensated to some extent if the biases are not shared.

1.12 (Max) Pooling

Typically pooling is used for down sampling, which means that the grid size of the data is reduced.

$$\begin{array}{c|c|c|c} 0 & 1 & 3 & 9 \\ \hline 4 & 8 & 6 & 2 \\ \hline 7 & 4 & 2 & 0 \\ \hline 4 & 3 & 5 & 1 \end{array} \mapsto \begin{array}{c|c} 8 & 9 \\ \hline 7 & 5 \end{array}$$

Example in 2-D: max pooling with 2×2 kernel and stride = 2.

Another benefit of max pooling with down sampling is that it is not so sensitive to small translations. If we take to grids, which are small translations of each other, and compare the pooled grid, we see that many of the values are same.

Pic: Without Shift, Shifted

1.13 Convolutional Network Components

Typically: (convolution, Relu, pooling) $\times N_1$, followed by FC $\times N_2$

1.14 Example of Learned Invariances

1.15 Examples of Architectures

We give here two examples of convolutional network architectures. Namely, one which is fairly classical and another which is from this year (2016).

...Explanations and pics...

1.16 Further Topics

In this last section we will mention briefly some of the topics presented in the Chapter 9 of the Deep Learning book that we have not yet covered.

There are variants of the usual way of doing the convolution...

*Comparison of Local Connections, Convolution, and FC (Optional)

*Tiled Convolution (optional)

*Recurrent Convolutional Networks

1.17 Links

* Our main source:

<http://www.deeplearningbook.org>

* Video with very basics of convolutional networks:

<https://www.youtube.com/watch?v=JiN9p5vWHDY>

* Video with slightly more advanced level introduction:

<https://www.youtube.com/watch?v=FmpDIaiMIeA>

* Video of a conference presentation:

<https://www.youtube.com/watch?v=AQirPKrAyDg>